

GRADED HOMEWORK 3 (Maximum Points: 87.5 points = 8.75% of the course grade)

Assigned: 10/4/2023

Group Portion (Max. 52.5 points) Due: 10/16/2023 at 01:30 PM on Canvas

Individual Portion (Max. 35 points) Due: 10/17/2023 at 11:59 PM on Canvas

NOTES:

- No late submission of either the group portion or the individual portion will be accepted.
- Read the submission instructions at the end of this document.
- If you do not work with your group to solve the Group Portion, you will get a zero score for the Group Portion. If you do not submit your Individual Portion, you will get a zero score for the entire homework regardless of the score you got on the Group Portion. Review the "Group Graded Homework Grading Policy" document and the "Graded Homework" section in the syllabus posted on Canvas.
- **ACADEMIC INTEGRITY:** The group portion must be done by your assigned group only; no collaboration with other groups or with anyone else is allowed. The individual portion must be done by you only; no collaboration with anyone (including your group members) is allowed. The use of generative AI tools (including ChatGPT, Bard, Bing Chat, and other AI writing and coding assistants) is not allowed in all the group questions and individual questions, except for those questions, if any, where the use of these tools is explicitly required. Violations of any of these rules will be considered academic misconduct and will result in action as specified in the Academic Integrity Code at The University of Oklahoma: <http://www.ou.edu/integrity>. Consult also the following web page for a Student's Guide to Academic Integrity at The University of Oklahoma: <http://www.ou.edu/integrity/students>.

PROBLEM 1:

Group Questions for Problem 1 (GQ1-GQ2):

GQ1) (8 points): To verify whether Azure SQL can check for integrity constraint violations, using the relational database that you have created for Problem 2 in Graded Homework 2, write SQL statements that implement the following queries, run them on Azure SQL, and capture the error messages generated by Azure SQL:

- a) One insertion query that violates the uniqueness of the primary key of a table.
- b) One insertion query that violates the not-null value of the primary key of a table.
- c) One insertion query, one deletion query, and one update query that violate the foreign key constraint of a table.
- d) One retrieval query that violates the domain constraint of an attribute of a table.

GQ2) (8 Points) Review the SQL file you have created for Problem 2 in Graded Homework 2, choose one table that should be indexed (assuming all queries are approximately equally frequent), write SQL statement(s) to create an index on that table, and rerun the queries that need to access the table and index. Provide your detailed explanations as to why you chose that table and that search key for indexing, whether that index is primary or secondary, and why you chose those queries to rerun.

Individual Question for Problem 1 (IQ1, 20 Points):

Each group member will be assigned one version of the following question randomly (**asking to consider only query A, B, C, or D being the most frequent**) after the group portion of the homework is due (see the submission instructions at the end of the document). You must be ready to answer any of the versions of the following question on your own to submit the individual portion of the homework:

Consider the case when **one of the queries** from the list below is executed far more frequently than the other queries from Problem 2 (GQ3) in Graded Homework 2:

- A) Find the age of the oldest performer who is either named “Hanks” or has acted in a movie named “The Departed”.
- B) Find the names of all movies that are either a Comedy or have had more than one performer act in them.
- C) Decrease the earnings of all directors who directed “Up” by 10%.
- D) Delete all movies released in the 70's and 80's ($1970 \leq \text{release_year} \leq 1989$).

Does this new information change your group's answer for the Group Question GQ2 of Problem 1 of this Graded Homework 3? Provide your detailed explanations as to why it does, or why it does not. If you argue for the change to the group's answer, make sure to explain which new table and/or which new search key should be used for indexing and whether the new index is primary or secondary. Provide the SQL statement(s) to create the new index. There's no need to execute the statement or any queries on Azure SQL.

PROBLEM 2:**Group Question for Problem 2 (CG3, 20 points):**

Write a **JAVA program** using JDBC and Azure SQL to implement the following options for the relational database that you have created for Problem 2 in Graded Homework 2:

1. Insert the pid, pname, (estimated) years_of_experience, and age of a new performer into table Performer. Assume that the value of years_of_experience is not known and will be estimated using the current information in the database as follows: the years_of_experience value should be set to the average of the years_of_experience values for all performers with an age within +/- 10 years of the new performer's age. If there are no such performers, the new performer's years_of_experience should be set to 18 less than his/her age value. After these calculations, the years_of_experience value should be adjusted so that it is at least 0 and no more than the performer's age value.
2. Insert the pid, pname, (estimated) years_of_experience, and age of a new performer into table Performer. Assume that the value of years_of_experience is not known and will be estimated using the current information in the database as follows: the years_of_experience value should be set to the average of the years_of_experience values for all performers who have acted in a movie that was directed by a director with a user-given did (director id). If there are no such performers, the new performer's years_of_experience should be set to 18 less than his/her age value. After these calculations, the years_of_experience value should be adjusted so that it is at least 0 and no more than the performer's age value.
3. Display the complete information of all performers.
4. Quit (exit the program).

Requirements:

- The program terminates only when the user chooses Option 4.
- Each of Options 1 and 2 must be implemented as a Transact SQL Stored Procedure and the pid, pname, age, and did (for option 2) must be entered as the procedure parameter values at runtime when the procedure

is called. No option logic must be handled by the Java program. It must only be used for user input collection, database call to execute the stored procedure, and the result presentation.

- For testing, execute Option 3 once before and after each execution of Options 1 and 2; execute Option 1 at least three times with different values of pid, pname, and age; execute Option 2 at least three times with different values of pid, pname, age, and user-given earnings value; and execute Option 4 at least one time to show that your program terminates correctly.
- The Java program and all the Transact SQL Stored Procedures must have proper in-line comments.

Individual Question for Problem 2 (IQ2, 20 points, every member will be assigned this question):

Modify your group's program for the Group Question GQ3 so that Query (2) is the following:

2. *Insert the pid, pname, (estimated) years_of_experience, and age of a new performer into table Performer. Assume that the value of years_of_experience is not known and will be estimated using the current information in the database as follows: the years_of_experience value should be set to the minimum of the years_of_experience values for all performers who have acted in a movie that was directed by a director whose earnings are above a user-given earnings value. If there are no such performers, the new performer's years_of_experience should be set to 18 less than his/her age value. After these calculations, the years_of_experience value should be adjusted so that it is at least 0 and no more than the performer's age value.*

Requirements:

- The program terminates only when the user chooses Option 4.
- Each of Options 1 and 2 must be implemented as a Transact SQL Stored Procedure and the pid, pname, age, and user-given earnings value (for Option 2) must be entered as the procedure parameter values at runtime when the procedure is called. No option logic must be handled by the Java program. It must only be used for user input collection, database call to execute the stored procedure, and the result presentation.
- For testing, execute Option 3 once before and after each execution of Options 1 and 2; execute Option 1 at least two times with different values of pid, pname, and age execute Option 2 at least two times with different values of pid, pname, age, and user-given earnings value; and execute Option 4 at least one time to show that your program terminates correctly.
- The Java program and all the Transact SQL Stored Procedures must have proper in-line comments.
- The Java program, all the Transact SQL Stored Procedures, and the testing must be implemented and executed using your own Azure SQL Database account.
- Solution Submission: Submit one PDF file that shows your SQL program, all the Transact SQL Store Procedures, the program execution for the required testing and its output.

PROBLEM 3 (Group Questions only, No Individual Question, No Azure SQL is required to solve this problem):

Given the following relational database table:

ComputerRepairSpecialists (name, certificateNo, city, price_per_hour)

The following 9 insertions are performed on the table *ComputerRepairSpecialists*:

*Insert record <Johnson, 11, Yukon, \$20>,
Insert record <Black, 33, OKC, \$20>,
Insert record <Grant, 22, Norman, \$15>,
Insert record <White, 77, OKC, \$20>,
Insert record <Chapman, 44, Edmond, \$20>,*

*Insert record <Ford, 66, Enid, \$25>,
Insert record <Haas, 99, OKC, \$20>,
Insert record <Hougen, 88, Yukon, \$25>,
Insert record <Clinton, 55, Tulsa, \$25>*

Assume each block in the *ComputerRepairSpecialists* file can store up to 3 specialist records and *ComputerRepairSpecialists* is organized as a sequential file with *name* as the ordering field.

Group Questions for Problem 3 (GQ4-GQ6):

GQ4) (5 Points): Show the contents (i.e. the data values as well as the associated block/bucket/record addresses) of the file after the last insertion.

GQ5) (5 Points): Assuming that *ComputerRepairSpecialists* is an index-sequential file on the search key *name*, show the contents of the primary index and the secondary index on *price_per_hour* after the last insertion.

GQ6) (6.5 Points): Assuming that a B-tree index file on *certificateNo* with order 3 is created for *ComputerRepairSpecialists* and no two specialists have the same *certificateNo*, using the definition of the B-tree given in class, show the content of the B-tree index file after the last insertion.

SUBMISSION INSTRUCTIONS:

Group Portion:

- The homework answers (one submission per group) must be submitted to **Canvas by 1:30 PM, 10/16/2023**. The submission includes the following:
 - Solutions for Problem 1: two files: a) the SQL file that shows the required explanations written as in-line comments, the SQL statement(s) for index creation, and the SQL queries that you chose to rerun; and b) The PDF file that shows the screenshots of the Azure SQL creation of the index and the Azure SQL execution of the SQL queries that you chose to rerun. Use the file name convention Group X_Problem1_HW3 where X is your group number.
 - Solutions for Problem 2: three files: a) a Java file (extension .java) containing the Java source program; b) an SQL file (extension .sql) containing the Transact SQL Stored Procedures; and c) a PDF file that shows the steps indicating that you have executed the program successfully (the output must be included). Use the file name convention Group X_Problem2_HW3 where X is your group number.
 - Solutions for Problem 3: A single PDF file with the typed/diagrammed solution for the Problem 3. No Azure SQL is required to solve Problem 3. Use the file name convention Group X_Problem3_HW3 where X is your group number.
 - Attach to your PDF for Problem 1 a cover page that contains the following information:

GROUP NUMBER: <write your group number here>

GROUP MEMBERS: <list the names of all members here>

GRADED HOMEWORK NUMBER: 3

COURSE: CS/DSA-4513 - DATABASE MANAGEMENT

SECTION: 001

SEMESTER: FALL 2023

INSTRUCTOR: LE GRUENWALD

SCORE: <<we will record the total score of your group for the problems 1, 2, and 3 here>>

Individual Portion:

- After the submission deadline of the group portion of this graded homework, and before the submission deadline of the individual portion of this graded homework, you will have to take a quiz on Canvas. The quiz will be open from 3:00 PM, Monday, October 16, 2023 to 11:59 PM, Tuesday, October 17, 2023. The quiz will contain the Individual Question IQ1 for Problem 1 and the Individual Question IQ2 for Problem 2. **Once you open the quiz, you will have 60 minutes to submit your answer. Only one attempt is allowed.** You will have to upload one PDF document as your answer. The quiz will also ask you for your feedback on your group members (i.e., the scores you give to each of your group members on the group portion of this graded homework) as outlined in the “Group Graded Homework Grading Policy” document available on Canvas. **If you do not submit your Individual Portion, you will receive a zero score for the entire homework regardless of the score you got on the Group Portion.**

NOTES:

- The instructions for “Using Java and JDBC to Connect to Azure SQL Database” and “Transact-SQL Stored Procedures” are available on Canvas.
- If you have questions concerning your Azure SQL account, using Java and JDBC to Connect to Azure SQL Database, Transact SQL Stored Procedures, or other questions concerning this homework, meet one of your TAs during their discussion hours or office hours over Zoom. Their discussion hours and office hours are on the Home Page on Canvas.
- **Start this project early to avoid last-minute system problems. No late submission will be accepted.**